

Potential Vulnerabilities in RIOT

Vuln 1: Out-of-Bound Write in Bluetooth Subsystem

Description

There is a potential OOB write in the `gen_prov_cont` function of the `pb_adv.c` file. This vulnerability occurs because the `seg` variable, which is user-controlled, is used to compute the offset where data is copied into. However, an integer underflow and overflow invalidates the validation on `seg` in [line 529](#).

Technical Details

1. `seg` is computed from `rx->gpc` in [line 483](#).
2. There is a validation in [line 529](#) that the offset computed from `seg` $((20 + ((seg - 1) * 23)))$ and the length of data does not overflow the receiver buffer.
3. However, if `seg = 0`, the computation will yield an offset of 253
 - a. `seg - 1` underflows, yielding 255
 - b. $255 * 23$ overflows multiple times, yielding 233
 - c. $233 + 20$ yields 253
4. If `buf->len` has a value greater than 2, then adding it to 253 overflows, yielding a number that is potentially less than `RX_BUFFER_MAX` (65).
5. Finally, the computed offset is used to compute the destination address of the `memcpy` in [line 534](#). This can cause a large out-of-bound write in the `rx.buf->data` array.

Potential Impact

An out-of-bound write can lead to an arbitrary code execution. This is more severe in real-time operating systems like Zephyr that run in embedded devices without common memory protection systems. Even on devices with some form of memory protection, this can still lead to a crash and a resultant denial of service.

Recommended Fix

Based on our analysis, we found that this vulnerability only occurs if `seg` has a value of 0. Hence, we suggest adding a validation that `seg` is not 0.

Vuln 2: Out-of-Bound Write in Bluetooth Subsystem

Description

There is a potential OOB Write vulnerability in the `gen_prov_start` function in `pb_adv.c`. The full length of the received data is copied into the `link.rx.buf` receiver buffer in [line 659](#) without any validation on the data size.

Technical Details

1. There is a `memcpy` on [line 659](#) that copies data into `link.rx.buf->data` buffer.
2. There is no validation on the length of the received data `buf->len`.
3. There seems to be a check on [line 617](#). However, `link.rx.buf->len` is read from the packet (in [line 604](#)) and may not correspond to the actual length of the packet in `buf->len`.
4. Hence, if the source buffer has a length greater than that of the receiver buffer, a malicious packet can write out of bounds in this function.

Potential Impact

An out-of-bound write can lead to an arbitrary code execution. This is more severe in real-time operating systems like Zephyr that run in embedded devices without common memory protection systems. Even on devices with some form of memory protection, this can still lead to a crash and a resultant denial of service.

Recommended Fix

We recommend that the validation on [line 617](#) should be updated so it validates that the `buf->len` field is not greater than the receiver buffer size.